

Homework 1

蔡云龙

21215068

Task 1 Frozen Lake MDP

1. (coding) Read through `vi_and_pi.py` and implement `policy_evaluation`, `policy_improvement` and `policy_iteration`. The stopping tolerance (defined as $\max_s |V_{old}(s) - V_{new}(s)|$) is $\text{tol} = 10^{-3}$. Use $\gamma = 0.9$. Return the optimal value function and the optimal policy.

策略迭代算法流程如下:

Policy Iteration (using iterative policy evaluation) for estimating $\pi \approx \pi_*$

1. Initialization
 $V(s) \in \mathbb{R}$ and $\pi(s) \in \mathcal{A}(s)$ arbitrarily for all $s \in \mathcal{S}$
2. Policy Evaluation
Loop:
 $\Delta \leftarrow 0$
Loop for each $s \in \mathcal{S}$:
 $v \leftarrow V(s)$
 $V(s) \leftarrow \sum_{s',r} p(s', r | s, \pi(s)) [r + \gamma V(s')]$
 $\Delta \leftarrow \max(\Delta, |v - V(s)|)$
until $\Delta < \theta$ (a small positive number determining the accuracy of estimation)
3. Policy Improvement
 $\text{policy-stable} \leftarrow \text{true}$
For each $s \in \mathcal{S}$:
 $\text{old-action} \leftarrow \pi(s)$
 $\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$
If $\text{old-action} \neq \pi(s)$, then $\text{policy-stable} \leftarrow \text{false}$
If policy-stable , then stop and return $V \approx v_*$ and $\pi \approx \pi_*$; else go to 2

<https://blog.csdn.net/njshake>

部分代码:

```
# policy_evaluation function
#####
# YOUR IMPLEMENTATION HERE #
while True:
    delta = 0.0
    for state in range(NS):
        value, action = value_function[state], policy[state]
        temp_val = 0
        for prob, next_state, reward, done in P[state][action]:
            temp_val += prob * (reward + gamma * value_function[next_state])
        value_function[state] = temp_val
        delta = max(delta, abs(value - value_function[state]))
    if delta < tol:
        break
#####

# policy_improvement function
#####
# YOUR IMPLEMENTATION HERE #
for state in range(NS):
    best_action, best_Q = None, -float('inf')
    for action in range(NA):
        current_Q = 0
        for prob, next_sate, reward, done in P[state][action]:
```



```
# value_iteration function
#####
# YOUR IMPLEMENTATION HERE #
for i in range(100000):
    delta = 0.0
    for state in range(nS):
        value = value_function[state]
        best_Q = -float('inf')

        for action in range(nA):
            this_Q = 0
            for (prob, next_state, reward, done) in P[state][action]:
                this_Q += prob * (reward + gamma * value_function[next_state])
            best_Q = max(best_Q, this_Q)
        value_function[state] = best_Q
        delta = max(delta, abs(value - value_function[state]))
    if delta < tol:
        break

policy = policy_improvement(P, nS, nA, value_function, policy, gamma)
#####
```

运行截图:

```
-----  
Beginning Value Iteration  
-----  
  
SFFF  
FHHH  
FFFF  
HFFG  
(Down)  
SFFF  
FHHH  
FFFF  
HFFG  
(Down)  
SFFF  
FHHH  
FFFF  
HFFG  
(Right)  
SFFF  
FHHH  
FFFF  
HFFG  
(Down)  
SFFF  
FHHH  
FFFF  
HFFG  
(Right)  
SFFF  
FHHH  
FFFF  
HFFG  
(Right)  
SFFF  
FHHH  
FFFF  
HFFG  
Episode reward: 1.000000
```

Task 2 Test Environment


1. **(written)** What is the maximum sum of rewards that can be achieved in a single trajectory in the test environment, assuming $\gamma = 1$? Show first that this value is attainable in a single trajectory, and then briefly argue why no other trajectory can achieve greater cumulative reward.

解: 状态过程为: $0 \rightarrow 2 \rightarrow 1 \rightarrow 2 \rightarrow 1 \rightarrow 0$, 即:

$$s_0 = 0, a_0 = 2, R_0 = 0.0, s_1 = 2, a_1 = 1, R_1 = 3.0, s_2 = 1, a_2 = 2, R_2 = 0.0, s_3 = 2, a_3 = 1, R_3 = 3.0, s_4 = 1, a_4 = 0, R_4 = 0.1, s_5 = 0$$

奖励为**6.1**。

图例如下:



说明：由表格可以看出最大奖励为从2到1，即3；每次执行2->1后要等一次才能再次执行2->1，又一共只有5步，可以执行两次奖励最高的步骤：2->1。一共奖励是6，最后只剩一步，需要从1到0，奖励为0.1，故总的最大奖励为6.1。过程如上所述。

Task 3 Tabular Q-Learning

1. (coding) Implement the `get_action` and `update` functions in `q_table.py`. Test your implementation by running `python q_table.py`.

`get_action` 函数中以 ϵ 的概率选择一个随机动作，这个过程通过 `env.action_space.sample()` 得到，否则返回参数 `best_action`，故有如下写法：

```

# get_action
#####
# YOUR IMPLEMENTATION HERE #
if np.random.uniform(0,1) < self.epsilon:
    return self.env.action_space.sample()
else:
    return best_action
#####

```

`update` 函数中根据当前步数 t 对 ϵ 进行线性更新，当 t 小于总步数时，线性从 `self.eps_begin` 变化到 `self.eps_end`，当 t 大于总步数时，保持为 `self.eps_end`：

```

# update
#####
# YOUR IMPLEMENTATION HERE #
if t <= self.nsteps:
    self.epsilon = self.eps_begin - t*(self.eps_begin-self.eps_end)/self.nsteps
else:
    self.epsilon = self.eps_end
#####

```

运行截图：

```

Test Results:
Test1: ok
Test2: ok
Test3: ok

```

Task 4 Maze Example

两种算法的流程和对比见下图：

Q-learning vs Sarsa

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

Initialize S ③

Choose A from S using policy derived from Q (e.g., ε -greedy)

Loop for each step of episode:

Take action A , observe R, S'

Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$ ① A' 为下一个state的实际action

$S \leftarrow S'; A \leftarrow A';$ ②

until S is terminal

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

Initialize S ③

Loop for each step of episode:

Choose A from S using policy derived from Q (e.g., ε -greedy)

Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$ ②

until S is terminal

① 默认 A' 为最优策略选的动作

仅 learn 部分代码不同，check_state_exist 与 choose_action 相同，如下：

```
# __init__
#####
# YOUR IMPLEMENTATION HERE #
self.q_table = pd.DataFrame(columns=self.actions, dtype=np.float64)
#####

# check_state_exist
#####
# YOUR IMPLEMENTATION HERE #
if state not in self.q_table.index:
    self.q_table = self.q_table.append(
        pd.Series(
            [0] * len(self.actions),
            index = self.q_table.columns,
            name = state,
        )
    )
#####

# choose_action
#####
# YOUR IMPLEMENTATION HERE #
self.check_state_exist(observation)
# action selection
if np.random.random() < self.epsilon:
    # choose best action
    state_action = self.q_table.loc[observation, :]

    action = np.random.choice(state_action[state_action == np.max(state_action)].index)
else:
    # choose random action
    action = np.random.choice(self.actions)
return action
#####
```

由于 SARSA 需要知道下一个动作但 Q Learning 不需要，并且在Q值的更新也不相同，因此两个类的 learn 函数不相同，分别如下：

1. (coding) Implement **Sarsa** in `RL_sarsa.py`.

```
def learn(self, state, action, reward, next_state, next_action):
    ''' update q table '''
    #####
    # YOUR IMPLEMENTATION HERE #
    self.check_state_exist(next_state)
    q_predict = self.q_table.loc[state, action]
    if next_state != 'terminal':
        q_target = reward + self.gamma * self.q_table.loc[next_state, next_action] # next state is not terminal
    else:
        q_target = reward # next state is terminal
    self.q_table.loc[state, action] += self.lr * (q_target - q_predict) # update
    #####
```

2. (coding) Implement **Q_learning** in `RL_q_learning.py`.

```
def learn(self, state, action, reward, next_state):
    ''' update q table '''
    #####
    # YOUR IMPLEMENTATION HERE #
    self.check_state_exist(next_state)
    q_predict = self.q_table.loc[state, action]
    if next_state != 'terminal':
        q_target = reward + self.gamma * self.q_table.loc[next_state, :].max() # next state is not terminal
    else:
        q_target = reward # next state is terminal
    self.q_table.loc[state, action] += self.lr * (q_target - q_predict) # update
    #####
```

3. 在 `run_this.py` 文件中，手动导入 `argparse` 库并添加了一个参数 `flag`，用来表示通过哪种方式运行，文件中的 `update` 函数和 `main` 函数如下(去注释)：

```
def update( flag ):
    for episode in range(100):
        # initial observation
        observation = env.reset()
        if flag == 'SARSA':
            action = RL.choose_action(str(observation))
        while True:
            env.render()

            #####
            # YOUR IMPLEMENTATION HERE #
            if flag == 'QLearning':
                action = RL.choose_action(str(observation))
            #####

            # RL take action and get next observation and reward
            observation_, reward, done = env.step(action)

            if flag == 'SARSA':
                action_ = RL.choose_action(str(observation_)) # next action for SARSA

            #####
            # YOUR IMPLEMENTATION HERE #
            if flag == 'SARSA':
                RL.learn(str(observation), action, reward, str(observation_), action_)
            elif flag == 'QLearning':
                RL.learn(str(observation), action, reward, str(observation_))
            #####

            # swap observation
            observation = observation_
```

```


        if flag == 'SARSA':
            action = action_
            # break while loop when end of this episode
            if done:
                break
        # end of game
        print('game over')
        env.destroy()

# main函数部分
if __name__ == "__main__":
    env = Maze()
    #####
    # YOUR IMPLEMENTATION HERE #
    flag = parser.parse_args().flag
    # flag = 1 # 1: SARSA 0: QLearning
    if flag == 'SARSA':
        print("Run with SRSA")
        RL = Sarsa(actions=list(range(env.n_actions)))
    elif flag == 'QLearning':
        print("Run with QLearning")
        RL = QLearning(actions=list(range(env.n_actions)))
    #####
    env.after(100, update(flag=flag))
    env.mainloop()

```

4. 运行时可以通过 `python run_this.py --flag QLearning` 或 `python run_this.py --flag SARSA` 来指定运行方式，效果如下：

PS E:\University\22-23研一下\强化学习\作业\Homework_1\RL_21215068_蔡云龙_homework1\task_3> `python run_this.py --flag SARSA`
Run with SRSA



PS E:\University\22-23研一下\强化学习\作业\Homework_1\RL_21215068_蔡云龙_homework1\task_3> `python run_this.py --flag QLearning`
Run with QLearning

